

# Building a GPT

Companion notebook to the [Zero To Hero](#) video on GPT.

```
# Kami selalu memulai dengan kumpulan data untuk dilatih. Mari unduh
kumpulan data shakespeare kecil
!wget https://raw.githubusercontent.com/karpathy/char-
rnn/master/data/tinyshakespeare/input.txt

--2025-01-04 15:11:34--
https://raw.githubusercontent.com/karpathy/char-rnn/master/data/tinysh
akespeare/input.txt
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|
185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 1115394 (1.1M) [text/plain]
Saving to: 'input.txt'

input.txt           0%[                  ] 0  ---KB/s
input.txt          100%[=====>] 1.06M  ---KB/s  in
0.04s

2025-01-04 15:11:34 (26.9 MB/s) - 'input.txt' saved [1115394/1115394]

# bacalah untuk memeriksanya
with open('input.txt', 'r', encoding='utf-8') as f:
    text = f.read()

print("length of dataset in characters: ", len(text))

length of dataset in characters: 1115394

# mari kita lihat 1000 karakter pertama
print(text[:1000])

First Citizen:
Before we proceed any further, hear me speak.

All:
Speak, speak.

First Citizen:
You are all resolved rather to die than to famish?

All:
Resolved. resolved.
```

First Citizen:  
First, you know Caius Marcius is chief enemy to the people.

All:  
We know't, we know't.

First Citizen:  
Let us kill him, and we'll have corn at our own price.  
Is't a verdict?

All:  
No more talking on't; let it be done: away, away!

Second Citizen:  
One word, good citizens.

First Citizen:  
We are accounted poor citizens, the patricians good.  
What authority surfeits on would relieve us: if they  
would yield us but the superfluity, while it were  
wholesome, we might guess they relieved us humanely;  
but they think we are too dear: the leanness that  
afflicts us, the object of our misery, is as an  
inventory to particularise their abundance; our  
sufferance is a gain to them Let us revenge this with  
our pikes, ere we become rakes: for the gods know I  
speak this in hunger for bread, not in thirst for revenge.

*# berikut adalah semua karakter unik yang muncul dalam teks ini*

```
chars = sorted(list(set(text)))  
vocab_size = len(chars)  
print(''.join(chars))  
print(vocab_size)
```

```
!$&',-.:;?ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz  
65
```

*# membuat pemetaan dari karakter ke bilangan bulat*

```
stoi = { ch:i for i,ch in enumerate(chars) }  
itos = { i:ch for i,ch in enumerate(chars) }  
encode = lambda s: [stoi[c] for c in s] # encoder: ambil string,  
keluarkan daftar bilangan bulat  
decode = lambda l: ''.join([itos[i] for i in l]) # decoder: ambil  
daftar bilangan bulat, keluarkan string
```

```
print(encode("hii there"))  
print(decode(encode("hii there")))
```

```
[46, 47, 47, 1, 58, 46, 43, 56, 43]
hii there
```

*# sekarang mari kita encode seluruh kumpulan data teks dan simpan ke dalam torch.Tensor*

```
import torch # kami menggunakan PyTorch: https://pytorch.org
```

```
data = torch.tensor(encode(text), dtype=torch.long)
```

```
print(data.shape, data.dtype)
```

*print(data[:1000]) # 1000 karakter yang kami lihat sebelumnya akan terlihat seperti ini di GPT*

```
torch.Size([1115394]) torch.int64
```

```
tensor([18, 47, 56, 57, 58,  1, 15, 47, 58, 47, 64, 43, 52, 10,  0,
        14, 43, 44,
         53, 56, 43,  1, 61, 43,  1, 54, 56, 53, 41, 43, 43, 42,  1,
        39, 52, 63,
         1, 44, 59, 56, 58, 46, 43, 56,  6,  1, 46, 43, 39, 56,  1,
        51, 43,  1,
         57, 54, 43, 39, 49,  8,  0,  0, 13, 50, 50, 10,  0, 31, 54,
        43, 39, 49,
         6,  1, 57, 54, 43, 39, 49,  8,  0,  0, 18, 47, 56, 57, 58,
        1, 15, 47,
         58, 47, 64, 43, 52, 10,  0, 37, 53, 59,  1, 39, 56, 43,  1,
        39, 50, 50,
         1, 56, 43, 57, 53, 50, 60, 43, 42,  1, 56, 39, 58, 46, 43,
        56,  1, 58,
         53,  1, 42, 47, 43,  1, 58, 46, 39, 52,  1, 58, 53,  1, 44,
        39, 51, 47,
         57, 46, 12,  0,  0, 13, 50, 50, 10,  0, 30, 43, 57, 53, 50,
        60, 43, 42,
         8,  1, 56, 43, 57, 53, 50, 60, 43, 42,  8,  0,  0, 18, 47,
        56, 57, 58,
         1, 15, 47, 58, 47, 64, 43, 52, 10,  0, 18, 47, 56, 57, 58,
        6,  1, 63,
         53, 59,  1, 49, 52, 53, 61,  1, 15, 39, 47, 59, 57,  1, 25,
        39, 56, 41,
         47, 59, 57,  1, 47, 57,  1, 41, 46, 47, 43, 44,  1, 43, 52,
        43, 51, 63,
         1, 58, 53,  1, 58, 46, 43,  1, 54, 43, 53, 54, 50, 43,  8,
        0,  0, 13,
         50, 50, 10,  0, 35, 43,  1, 49, 52, 53, 61,  5, 58,  6,  1,
        61, 43,  1,
         49, 52, 53, 61,  5, 58,  8,  0,  0, 18, 47, 56, 57, 58,  1,
        15, 47, 58,
         47, 64, 43, 52, 10,  0, 24, 43, 58,  1, 59, 57,  1, 49, 47,
        50, 50,  1,
         46, 47, 51,  6,  1, 39, 52, 42,  1, 61, 43,  5, 50, 50,  1,
        46, 39, 60,
         43,  1, 41, 53, 56, 52,  1, 39, 58,  1, 53, 59, 56,  1, 53,
        61, 52,  1,
```

	54, 56, 47, 41, 43, 8, 0, 21, 57, 5, 58, 1, 39, 1, 60,
43, 56,	42,
	47, 41, 58, 12, 0, 0, 13, 50, 50, 10, 0, 26, 53, 1, 51,
53, 56,	43,
	1, 58, 39, 50, 49, 47, 52, 45, 1, 53, 52, 5, 58, 11, 1,
50, 43,	58,
	1, 47, 58, 1, 40, 43, 1, 42, 53, 52, 43, 10, 1, 39, 61,
39, 63,	6,
	1, 39, 61, 39, 63, 2, 0, 0, 31, 43, 41, 53, 52, 42, 1,
15, 47,	58,
	47, 64, 43, 52, 10, 0, 27, 52, 43, 1, 61, 53, 56, 42, 6,
1, 45,	53,
	53, 42, 1, 41, 47, 58, 47, 64, 43, 52, 57, 8, 0, 0, 18,
47, 56,	57,
	58, 1, 15, 47, 58, 47, 64, 43, 52, 10, 0, 35, 43, 1, 39,
56, 43,	1,
	39, 41, 41, 53, 59, 52, 58, 43, 42, 1, 54, 53, 53, 56, 1,
41, 47,	58,
	47, 64, 43, 52, 57, 6, 1, 58, 46, 43, 1, 54, 39, 58, 56,
47, 41,	47,
	39, 52, 57, 1, 45, 53, 53, 42, 8, 0, 35, 46, 39, 58, 1,
39, 59,	58,
	46, 53, 56, 47, 58, 63, 1, 57, 59, 56, 44, 43, 47, 58, 57,
1, 53,	52,
	1, 61, 53, 59, 50, 42, 1, 56, 43, 50, 47, 43, 60, 43, 1,
59, 57,	10,
	1, 47, 44, 1, 58, 46, 43, 63, 0, 61, 53, 59, 50, 42, 1,
63, 47,	43,
	50, 42, 1, 59, 57, 1, 40, 59, 58, 1, 58, 46, 43, 1, 57,
59, 54,	43,
	56, 44, 50, 59, 47, 58, 63, 6, 1, 61, 46, 47, 50, 43, 1,
47, 58,	1,
	61, 43, 56, 43, 0, 61, 46, 53, 50, 43, 57, 53, 51, 43, 6,
1, 61,	43,
	1, 51, 47, 45, 46, 58, 1, 45, 59, 43, 57, 57, 1, 58, 46,
43, 63,	1,
	56, 43, 50, 47, 43, 60, 43, 42, 1, 59, 57, 1, 46, 59, 51,
39, 52,	43,
	50, 63, 11, 0, 40, 59, 58, 1, 58, 46, 43, 63, 1, 58, 46,
47, 52,	49,
	1, 61, 43, 1, 39, 56, 43, 1, 58, 53, 53, 1, 42, 43, 39,
56, 10,	1,
	58, 46, 43, 1, 50, 43, 39, 52, 52, 43, 57, 57, 1, 58, 46,
39, 58,	0,
	39, 44, 44, 50, 47, 41, 58, 57, 1, 59, 57, 6, 1, 58, 46,
43, 1,	53,
	40, 48, 43, 41, 58, 1, 53, 44, 1, 53, 59, 56, 1, 51, 47,
57, 43,	56,
	63, 6, 1, 47, 57, 1, 39, 57, 1, 39, 52, 0, 47, 52, 60,

```

43, 52, 58,
53, 56, 63, 1, 58, 53, 1, 54, 39, 56, 58, 47, 41, 59, 50,
39, 56, 47,
57, 43, 1, 58, 46, 43, 47, 56, 1, 39, 40, 59, 52, 42, 39,
52, 41, 43,
11, 1, 53, 59, 56, 0, 57, 59, 44, 44, 43, 56, 39, 52, 41,
43, 1, 47,
57, 1, 39, 1, 45, 39, 47, 52, 1, 58, 53, 1, 58, 46, 43,
51, 1, 24,
43, 58, 1, 59, 57, 1, 56, 43, 60, 43, 52, 45, 43, 1, 58,
46, 47, 57,
1, 61, 47, 58, 46, 0, 53, 59, 56, 1, 54, 47, 49, 43, 57,
6, 1, 43,
56, 43, 1, 61, 43, 1, 40, 43, 41, 53, 51, 43, 1, 56, 39,
49, 43, 57,
10, 1, 44, 53, 56, 1, 58, 46, 43, 1, 45, 53, 42, 57, 1,
49, 52, 53,
61, 1, 21, 0, 57, 54, 43, 39, 49, 1, 58, 46, 47, 57, 1,
47, 52, 1,
46, 59, 52, 45, 43, 56, 1, 44, 53, 56, 1, 40, 56, 43, 39,
42, 6, 1,
52, 53, 58, 1, 47, 52, 1, 58, 46, 47, 56, 57, 58, 1, 44,
53, 56, 1,
56, 43, 60, 43, 52, 45, 43, 8, 0, 0])

```

*# Sekarang mari kita bagi data menjadi rangkaian pelatihan dan validasi*

```

n = int(0.9*len(data)) # 90% pertama akan dilatih, sisanya val
train_data = data[:n]
val_data = data[n:]

```

```

block_size = 8
train_data[:block_size+1]

tensor([18, 47, 56, 57, 58, 1, 15, 47, 58])

```

```

x = train_data[:block_size]
y = train_data[1:block_size+1]
for t in range(block_size):
    context = x[:t+1]
    target = y[t]
    print(f"when input is {context} the target: {target}")

```

```

when input is tensor([18]) the target: 47
when input is tensor([18, 47]) the target: 56
when input is tensor([18, 47, 56]) the target: 57
when input is tensor([18, 47, 56, 57]) the target: 58
when input is tensor([18, 47, 56, 57, 58]) the target: 1
when input is tensor([18, 47, 56, 57, 58, 1]) the target: 15

```

when input is tensor([18, 47, 56, 57, 58, 1, 15]) the target: 47  
when input is tensor([18, 47, 56, 57, 58, 1, 15, 47]) the target: 58

```
torch.manual_seed(1337)
batch_size = 4 # berapa banyak barisan independen yang akan kita
proses secara paralel?
block_size = 8 # berapa panjang konteks maksimum untuk prediksi?
```

```
def get_batch(split):
    # menghasilkan sejumlah kecil input data x dan target y
    data = train_data if split == 'train' else val_data
    ix = torch.randint(len(data) - block_size, (batch_size,))
    x = torch.stack([data[i:i+block_size] for i in ix])
    y = torch.stack([data[i+1:i+block_size+1] for i in ix])
    return x, y
```

```
xb, yb = get_batch('train')
print('inputs:')
print(xb.shape)
print(xb)
print('targets:')
print(yb.shape)
print(yb)
```

```
print('----')
```

```
for b in range(batch_size): # dimensi kumpulan
    for t in range(block_size): # dimensi waktu
        context = xb[b, :t+1]
        target = yb[b,t]
        print(f"when input is {context.tolist()} the target:
{target}")
```

```
inputs:
torch.Size([4, 8])
tensor([[24, 43, 58, 5, 57, 1, 46, 43],
        [44, 53, 56, 1, 58, 46, 39, 58],
        [52, 58, 1, 58, 46, 39, 58, 1],
        [25, 17, 27, 10, 0, 21, 1, 54]])
```

```
targets:
torch.Size([4, 8])
tensor([[43, 58, 5, 57, 1, 46, 43, 39],
        [53, 56, 1, 58, 46, 39, 58, 1],
        [58, 1, 58, 46, 39, 58, 1, 46],
        [17, 27, 10, 0, 21, 1, 54, 39]])
```

```
----
```

```
when input is [24] the target: 43
when input is [24, 43] the target: 58
when input is [24, 43, 58] the target: 5
when input is [24, 43, 58, 5] the target: 57
```

```

when input is [24, 43, 58, 5, 57] the target: 1
when input is [24, 43, 58, 5, 57, 1] the target: 46
when input is [24, 43, 58, 5, 57, 1, 46] the target: 43
when input is [24, 43, 58, 5, 57, 1, 46, 43] the target: 39
when input is [44] the target: 53
when input is [44, 53] the target: 56
when input is [44, 53, 56] the target: 1
when input is [44, 53, 56, 1] the target: 58
when input is [44, 53, 56, 1, 58] the target: 46
when input is [44, 53, 56, 1, 58, 46] the target: 39
when input is [44, 53, 56, 1, 58, 46, 39] the target: 58
when input is [44, 53, 56, 1, 58, 46, 39, 58] the target: 1
when input is [52] the target: 58
when input is [52, 58] the target: 1
when input is [52, 58, 1] the target: 58
when input is [52, 58, 1, 58] the target: 46
when input is [52, 58, 1, 58, 46] the target: 39
when input is [52, 58, 1, 58, 46, 39] the target: 58
when input is [52, 58, 1, 58, 46, 39, 58] the target: 1
when input is [52, 58, 1, 58, 46, 39, 58, 1] the target: 46
when input is [25] the target: 17
when input is [25, 17] the target: 27
when input is [25, 17, 27] the target: 10
when input is [25, 17, 27, 10] the target: 0
when input is [25, 17, 27, 10, 0] the target: 21
when input is [25, 17, 27, 10, 0, 21] the target: 1
when input is [25, 17, 27, 10, 0, 21, 1] the target: 54
when input is [25, 17, 27, 10, 0, 21, 1, 54] the target: 39

```

```

print(xb) # masukan kita ke trafo

```

```

tensor([[24, 43, 58, 5, 57, 1, 46, 43],
        [44, 53, 56, 1, 58, 46, 39, 58],
        [52, 58, 1, 58, 46, 39, 58, 1],
        [25, 17, 27, 10, 0, 21, 1, 54]])

```

```

import torch
import torch.nn as nn
from torch.nn import functional as F
torch.manual_seed(1337)

```

```

class BigramLanguageModel(nn.Module):

```

```

    def __init__(self, vocab_size):
        super().__init__()
        # setiap token secara langsung membaca logit untuk token
berikutnya dari tabel pencarian
        self.token_embedding_table = nn.Embedding(vocab_size,
        vocab_size)

```

```

def forward(self, idx, targets=None):
    # idx dan target keduanya merupakan tensor bilangan bulat (B,T).
    logits = self.token_embedding_table(idx) # (B,T,C)

    if targets is None:
        loss = None
    else:
        B, T, C = logits.shape
        logits = logits.view(B*T, C)
        targets = targets.view(B*T)
        loss = F.cross_entropy(logits, targets)

    return logits, loss

def generate(self, idx, max_new_tokens):
    # idx adalah array indeks (B, T) dalam konteks saat ini
    for _ in range(max_new_tokens):
        # dapatkan prediksinya
        logits, loss = self(idx)
        # fokus hanya pada langkah terakhir kali
        logits = logits[:, -1, :] # menjadi (B,C)
        # terapkan softmax untuk mendapatkan probabilitas
        probs = F.softmax(logits, dim=-1) # (B, C)
        # sampel dari distribusi
        idx_next = torch.multinomial(probs, num_samples=1) # (B, 1)

        # tambahkan indeks sampel ke urutan yang sedang berjalan
        idx = torch.cat((idx, idx_next), dim=1) # (B, T+1)
    return idx

m = BigramLanguageModel(vocab_size)
logits, loss = m(xb, yb)
print(logits.shape)
print(loss)

print(decode(m.generate(idx = torch.zeros((1, 1), dtype=torch.long),
max_new_tokens=100)[0].tolist()))

torch.Size([32, 65])
tensor(4.8786, grad_fn=<NllLossBackward0>)

Sr?qP-QWktXoL&jLDJg0LVz'RIoDqHdhsV&vLLxatjscMpwLERSPyao.qfzs$Ys$zF-
w,;eEkzxjgCKFChs!iWW.ObzDnxA Ms$3

# buat pengoptimal PyTorch
optimizer = torch.optim.AdamW(m.parameters(), lr=1e-3)

batch_size = 32
for steps in range(100): # tingkatan jumlah langkah untuk hasil yang

```



baik...

```
# sampel kumpulan data
xb, yb = get_batch('train')

# evaluasi kerugian
logits, loss = m(xb, yb)
optimizer.zero_grad(set_to_none=True)
loss.backward()
optimizer.step()
```

```
print(loss.item())
```

4.587916374206543

```
print(decode(m.generate(idx = torch.zeros((1, 1), dtype=torch.long),
max_new_tokens=500)[0].tolist()))
```

xiKi-RJ:CgqVuUa!U?qMH.uk!sCuMXvv!CJFfx;LgRyJkn0Eti.?I&-gPLLyulId?  
XlaInQ'q,lT\$  
3Q&sGlvHQ?mqSq-eON  
x?SP fUAFCAuCX:bOlgiRQWN:Mphaw  
tRLKuYXEaAXxrcq-gCUzeh3w!AcyaylgYWjmJM?  
Uzw:inaY,:C&0ECW:vmGGJAn3onAuMgia!ms\$Vb q-gC0cPcUh0nxJGUGSPJWT:..?  
ujmJFoiNL&A'DxY,prZ?qdT;hoo'dHooXXlxf'WkHK&u3Q?rqUi.kz;?Yx?  
C&u3Qbfzxlyh'Vl:zyxjKXgC?  
lv'QKFiBeviNx0'm!Upm\$srm&TqViqiBD3HBP!juE0pmZJyF\$Fwfy!PlvWPFC  
&WdDP!Ko,px  
x  
tRE0E;AJ.BeXkyLOVD3KHps\$e?nD,.SFbWWI'ubcL!q-tU;aXmJ&uGXHxJXI&Z!  
gHRpajj;l.  
pTErIBjx;JKIgoCnLGXrJSP!AU-AcbczR?

## The mathematical trick in self-attention

```
# contoh mainan yang mengilustrasikan bagaimana perkalian matriks
dapat digunakan untuk "agregasi berbobot"
torch.manual_seed(42)
a = torch.tril(torch.ones(3, 3))
a = a / torch.sum(a, 1, keepdim=True)
b = torch.randint(0, 10, (3, 2)).float()
c = a @ b
print('a=')
print(a)
print('--')
print('b=')
print(b)
print('--')
```

```

print('c=')
print(c)

a=
tensor([[1.0000, 0.0000, 0.0000],
        [0.5000, 0.5000, 0.0000],
        [0.3333, 0.3333, 0.3333]])
--
b=
tensor([[2., 7.],
        [6., 4.],
        [6., 5.]])
--
c=
tensor([[2.0000, 7.0000],
        [4.0000, 5.5000],
        [4.6667, 5.3333]])

# perhatikan contoh mainan berikut :

torch.manual_seed(1337)
B,T,C = 4,8,2 # batch, waktu, saluran
x = torch.randn(B,T,C)
x.shape

torch.Size([4, 8, 2])

# Kita ingin  $x[b,t] = \text{mean}_{i \leq t} x[b,i]$ 
xbow = torch.zeros((B,T,C))
for b in range(B):
    for t in range(T):
        xprev = x[b,:t+1] # (t,C)
        xbow[b,t] = torch.mean(xprev, 0)

# versi 2: menggunakan perkalian matriks untuk agregasi berbobot
wei = torch.tril(torch.ones(T, T))
wei = wei / wei.sum(1, keepdim=True)
xbow2 = wei @ x # (B, T, T) @ (B, T, C) ----> (B, T, C)
torch.allclose(xbow, xbow2)

False

# versi 3: gunakan Softmax
tril = torch.tril(torch.ones(T, T))
wei = torch.zeros((T,T))
wei = wei.masked_fill(tril == 0, float('-inf'))
wei = F.softmax(wei, dim=-1)
xbow3 = wei @ x
torch.allclose(xbow, xbow3)

False

```

```

# versi 4: perhatian diri!
torch.manual_seed(1337)
B,T,C = 4,8,32 # batch, waktu, saluran
x = torch.randn(B,T,C)

# mari kita lihat seorang Kepala melakukan perhatian diri
head_size = 16
key = nn.Linear(C, head_size, bias=False)
query = nn.Linear(C, head_size, bias=False)
value = nn.Linear(C, head_size, bias=False)
k = key(x) # (B, T, 16)
q = query(x) # (B, T, 16)
wei = q @ k.transpose(-2, -1) # (B, T, 16) @ (B, 16, T) ---> (B, T, T)

tril = torch.tril(torch.ones(T, T))
#wei = torch.zeros((T,T))
wei = wei.masked_fill(tril == 0, float('-inf'))
wei = F.softmax(wei, dim=-1)

v = value(x)
out = wei @ v
#out = wei @ x

out.shape
torch.Size([4, 8, 16])

wei[0]
tensor([[1.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000],
        [0.1574, 0.8426, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000],
        [0.2088, 0.1646, 0.6266, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000],
        [0.5792, 0.1187, 0.1889, 0.1131, 0.0000, 0.0000, 0.0000, 0.0000],
        [0.0294, 0.1052, 0.0469, 0.0276, 0.7909, 0.0000, 0.0000, 0.0000],
        [0.0176, 0.2689, 0.0215, 0.0089, 0.6812, 0.0019, 0.0000, 0.0000],
        [0.1691, 0.4066, 0.0438, 0.0416, 0.1048, 0.2012, 0.0329, 0.0000],
        [0.0210, 0.0843, 0.0555, 0.2297, 0.0573, 0.0709, 0.2423, 0.2391]]),
      grad_fn=<SelectBackward0>)
```

Notes:

- Attention is a **communication mechanism**. Can be seen as nodes in a directed graph looking at each other and aggregating information with a weighted sum from all nodes that point to them, with data-dependent weights.
- There is no notion of space. Attention simply acts over a set of vectors. This is why we need to positionally encode tokens.
- Each example across batch dimension is of course processed completely independently and never "talk" to each other
- In an "encoder" attention block just delete the single line that does masking with `tril`, allowing all tokens to communicate. This block here is called a "decoder" attention block because it has triangular masking, and is usually used in autoregressive settings, like language modeling.
- "self-attention" just means that the keys and values are produced from the same source as queries. In "cross-attention", the queries still get produced from  $x$ , but the keys and values come from some other, external source (e.g. an encoder module)
- "Scaled" attention additionally divides `wei` by  $1/\sqrt{\text{head\_size}}$ . This makes it so when input  $Q, K$  are unit variance, `wei` will be unit variance too and Softmax will stay diffuse and not saturate too much. Illustration below

```
k = torch.randn(B,T,head_size)
q = torch.randn(B,T,head_size)
wei = q @ k.transpose(-2, -1) * head_size**-0.5

k.var()
tensor(1.0449)

q.var()
tensor(1.0700)

wei.var()
tensor(1.0918)

torch.softmax(torch.tensor([0.1, -0.2, 0.3, -0.2, 0.5]), dim=-1)
tensor([0.1925, 0.1426, 0.2351, 0.1426, 0.2872])

torch.softmax(torch.tensor([0.1, -0.2, 0.3, -0.2, 0.5])*8, dim=-1) #
menjadi terlalu pucat, menyatu menjadi satu-panas
tensor([0.0326, 0.0030, 0.1615, 0.0030, 0.8000])

class LayerNorm1d: # (dulu BatchNorm1d)

    def __init__(self, dim, eps=1e-5, momentum=0.1):
        self.eps = eps
        self.gamma = torch.ones(dim)
        self.beta = torch.zeros(dim)

    def __call__(self, x):
```

```

    # hitung umpan maju
    xmean = x.mean(1, keepdim=True) # maksud kumpulan
    xvar = x.var(1, keepdim=True) # varians batch
    xhat = (x - xmean) / torch.sqrt(xvar + self.eps) # normalisasi ke
unit varians
    self.out = self.gamma * xhat + self.beta
    return self.out

def parameters(self):
    return [self.gamma, self.beta]

torch.manual_seed(1337)
module = LayerNorm1d(100)
x = torch.randn(32, 100) # ukuran batch 32 dari vektor 100 dimensi
x = module(x)
x.shape

torch.Size([32, 100])

x[:,0].mean(), x[:,0].std() # mean,std dari satu fitur di semua input
batch

(tensor(0.1469), tensor(0.8803))

x[0,:].mean(), x[0,:].std() # mean,std dari satu masukan dari batch,
fitur-fiturnya

(tensor(-9.5367e-09), tensor(1.0000))

# Contoh terjemahan Bahasa Prancis ke Bahasa Inggris:

# <----- ENCODE -----><----- DECODE
----->
# les réseaux de neurones sont géniaux! <START> jaringan neural luar
biasa!<END>

```

## Full finished code, for reference

You may want to refer directly to the git repo instead though.

```

import torch
import torch.nn as nn
from torch.nn import functional as F

# hyperparameters
batch_size = 16 #berapa banyak barisan independen yang akan kita
proses secara paralel?
block_size = 32 # berapa panjang konteks maksimum untuk prediksi?
max_iters = 5000
eval_interval = 100

```

```

learning_rate = 1e-3
device = 'cuda' if torch.cuda.is_available() else 'cpu'
eval_iters = 200
n_embd = 64
n_head = 4
n_layer = 4
dropout = 0.0
# -----

torch.manual_seed(1337)

# wget
https://raw.githubusercontent.com/karpathy/char-rnn/master/data/tinyshakespeare/input.txt
with open('input.txt', 'r', encoding='utf-8') as f:
    text = f.read()

# berikut adalah semua karakter unik yang muncul dalam teks ini
chars = sorted(list(set(text)))
vocab_size = len(chars)
# membuat pemetaan dari karakter ke bilangan bulat
stoi = { ch:i for i,ch in enumerate(chars) }
itos = { i:ch for i,ch in enumerate(chars) }
encode = lambda s: [stoi[c] for c in s] # encoder: ambil string,
keluarkan daftar bilangan bulat
decode = lambda l: ''.join([itos[i] for i in l]) # decoder: ambil
daftar bilangan bulat, keluarkan string

# Latih dan uji perpecahan
data = torch.tensor(encode(text), dtype=torch.long)
n = int(0.9*len(data)) # 90% pertama akan dilatih, sisanya val
train_data = data[:n]
val_data = data[n:]

# memuat data
def get_batch(split):
    # menghasilkan sejumlah kecil data input x dan target y
    data = train_data if split == 'train' else val_data
    ix = torch.randint(len(data) - block_size, (batch_size,))
    x = torch.stack([data[i:i+block_size] for i in ix])
    y = torch.stack([data[i+1:i+block_size+1] for i in ix])
    x, y = x.to(device), y.to(device)
    return x, y

@torch.no_grad()
def estimate_loss():
    out = {}
    model.eval()
    for split in ['train', 'val']:
        losses = torch.zeros(eval_iters)

```

```

        for k in range(eval_iters):
            X, Y = get_batch(split)
            logits, loss = model(X, Y)
            losses[k] = loss.item()
        out[split] = losses.mean()
    model.train()
    return out

class Head(nn.Module):
    """ one head of self-attention """

    def __init__(self, head_size):
        super().__init__()
        self.key = nn.Linear(n_embd, head_size, bias=False)
        self.query = nn.Linear(n_embd, head_size, bias=False)
        self.value = nn.Linear(n_embd, head_size, bias=False)
        self.register_buffer('tril', torch.tril(torch.ones(block_size,
block_size)))

        self.dropout = nn.Dropout(dropout)

    def forward(self, x):
        B, T, C = x.shape
        k = self.key(x) # (B, T, C)
        q = self.query(x) # (B, T, C)
        # menghitung skor perhatian ("afinitas")
        wei = q @ k.transpose(-2, -1) * C**-0.5 # (B, T, C) @ (B, C, T)
        -> (B, T, T)
        wei = wei.masked_fill(self.tril[:T, :T] == 0, float('-inf')) #
(B, T, T)
        wei = F.softmax(wei, dim=-1) # (B, T, T)
        wei = self.dropout(wei)
        # melakukan agregasi nilai tertimbang
        v = self.value(x) # (B, T, C)
        out = wei @ v # (B, T, T) @ (B, T, C) -> (B, T, C)
        return out

class MultiHeadAttention(nn.Module):
    """ multiple heads of self-attention in parallel """

    def __init__(self, num_heads, head_size):
        super().__init__()
        self.heads = nn.ModuleList([Head(head_size) for _ in
range(num_heads)])
        self.proj = nn.Linear(n_embd, n_embd)
        self.dropout = nn.Dropout(dropout)

    def forward(self, x):
        out = torch.cat([h(x) for h in self.heads], dim=-1)
        out = self.dropout(self.proj(out))

```

```

        return out

class FeedFoward(nn.Module):
    """ a simple linear layer followed by a non-linearity """

    def __init__(self, n_embd):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(n_embd, 4 * n_embd),
            nn.ReLU(),
            nn.Linear(4 * n_embd, n_embd),
            nn.Dropout(dropout),
        )

    def forward(self, x):
        return self.net(x)

class Block(nn.Module):
    """ Transformer block: communication followed by computation """

    def __init__(self, n_embd, n_head):
        # n_embd: menyematkan dimensi, n_head: jumlah kepala yang kita
        # inginkan
        super().__init__()
        head_size = n_embd // n_head
        self.sa = MultiHeadAttention(n_head, head_size)
        self.ffwd = FeedFoward(n_embd)
        self.ln1 = nn.LayerNorm(n_embd)
        self.ln2 = nn.LayerNorm(n_embd)

    def forward(self, x):
        x = x + self.sa(self.ln1(x))
        x = x + self.ffwd(self.ln2(x))
        return x

# model bigram super sederhana
class BigramLanguageModel(nn.Module):

    def __init__(self):
        super().__init__()
        # setiap token secara langsung membaca logit untuk token
        # berikutnya dari tabel pencarian
        self.token_embedding_table = nn.Embedding(vocab_size, n_embd)
        self.position_embedding_table = nn.Embedding(block_size,
n_embd)
        self.blocks = nn.Sequential(*[Block(n_embd, n_head=n_head) for
_ in range(n_layer)])
        self.ln_f = nn.LayerNorm(n_embd) # norma lapisan akhir
        self.lm_head = nn.Linear(n_embd, vocab_size)

```



```

def forward(self, idx, targets=None):
    B, T = idx.shape

    # idx dan target keduanya merupakan tensor bilangan bulat
    (B,T).
    tok_emb = self.token_embedding_table(idx) # (B,T,C)
    pos_emb = self.position_embedding_table(torch.arange(T,
device=device)) # (T,C)
    x = tok_emb + pos_emb # (B,T,C)
    x = self.blocks(x) # (B,T,C)
    x = self.ln_f(x) # (B,T,C)
    logits = self.lm_head(x) # (B,T,vocab_size)

    if targets is None:
        loss = None
    else:
        B, T, C = logits.shape
        logits = logits.view(B*T, C)
        targets = targets.view(B*T)
        loss = F.cross_entropy(logits, targets)

    return logits, loss

def generate(self, idx, max_new_tokens):
    # idx adalah array indeks (B, T) dalam konteks saat ini
    for _ in range(max_new_tokens):
        # pangkas idx ke token block_size terakhir
        idx_cond = idx[:, -block_size:]
        # dapatkan prediksinya
        logits, loss = self(idx_cond)
        # fokus hanya pada langkah terakhir kali
        logits = logits[:, -1, :] # menjadi (B, C)
        # terapkan softmax untuk mendapatkan probabilitas
        probs = F.softmax(logits, dim=-1) # (B, C)
        # sampel dari distribusi
        idx_next = torch.multinomial(probs, num_samples=1) # (B,
1)
        # tambahkan indeks sampel ke urutan yang sedang berjalan
        idx = torch.cat((idx, idx_next), dim=1) # (B, T+1)
    return idx

model = BigramLanguageModel()
m = model.to(device)
# mencetak jumlah parameter dalam model
print(sum(p.numel() for p in m.parameters())/1e6, 'M parameters')

# buat pengoptimal PyTorch
optimizer = torch.optim.AdamW(model.parameters(), lr=learning_rate)

for iter in range(max_iters):

```

```

# sesekali mengevaluasi kerugian pada set kereta dan val
if iter % eval_interval == 0 or iter == max_iters - 1:
    losses = estimate_loss()
    print(f"step {iter}: train loss {losses['train']:.4f}, val
loss {losses['val']:.4f}")

# sampel kumpulan data
xb, yb = get_batch('train')

# evaluasi kerugian
logits, loss = model(xb, yb)
optimizer.zero_grad(set_to_none=True)
loss.backward()
optimizer.step()

# hasilkan dari model
context = torch.zeros((1, 1), dtype=torch.long, device=device)
print(decode(m.generate(context, max_new_tokens=2000)[0].tolist()))

0.209729 M parameters
step 0: train loss 4.4116, val loss 4.4022
step 100: train loss 2.6568, val loss 2.6670
step 200: train loss 2.5091, val loss 2.5058
step 300: train loss 2.4197, val loss 2.4336
step 400: train loss 2.3501, val loss 2.3562
step 500: train loss 2.2963, val loss 2.3125
step 600: train loss 2.2407, val loss 2.2496
step 700: train loss 2.2054, val loss 2.2187
step 800: train loss 2.1633, val loss 2.1866
step 900: train loss 2.1241, val loss 2.1504
step 1000: train loss 2.1036, val loss 2.1306
step 1100: train loss 2.0698, val loss 2.1180
step 1200: train loss 2.0380, val loss 2.0791
step 1300: train loss 2.0248, val loss 2.0634
step 1400: train loss 1.9926, val loss 2.0359
step 1500: train loss 1.9697, val loss 2.0287
step 1600: train loss 1.9627, val loss 2.0477
step 1700: train loss 1.9403, val loss 2.0115
step 1800: train loss 1.9090, val loss 1.9941
step 1900: train loss 1.9092, val loss 1.9858
step 2000: train loss 1.8847, val loss 1.9925
step 2100: train loss 1.8724, val loss 1.9757
step 2200: train loss 1.8580, val loss 1.9594
step 2300: train loss 1.8560, val loss 1.9537
step 2400: train loss 1.8412, val loss 1.9427
step 2500: train loss 1.8141, val loss 1.9402
step 2600: train loss 1.8292, val loss 1.9397
step 2700: train loss 1.8116, val loss 1.9322
step 2800: train loss 1.8032, val loss 1.9218

```

step 2900: train loss 1.8022, val loss 1.9285  
step 3000: train loss 1.7955, val loss 1.9195  
step 3100: train loss 1.7672, val loss 1.9192  
step 3200: train loss 1.7568, val loss 1.9138  
step 3300: train loss 1.7551, val loss 1.9059  
step 3400: train loss 1.7549, val loss 1.8945  
step 3500: train loss 1.7383, val loss 1.8956  
step 3600: train loss 1.7242, val loss 1.8868  
step 3700: train loss 1.7273, val loss 1.8822  
step 3800: train loss 1.7176, val loss 1.8923  
step 3900: train loss 1.7219, val loss 1.8750  
step 4000: train loss 1.7131, val loss 1.8603  
step 4100: train loss 1.7105, val loss 1.8777  
step 4200: train loss 1.7033, val loss 1.8675  
step 4300: train loss 1.7038, val loss 1.8556  
step 4400: train loss 1.7057, val loss 1.8643  
step 4500: train loss 1.6875, val loss 1.8528  
step 4600: train loss 1.6887, val loss 1.8405  
step 4700: train loss 1.6834, val loss 1.8501  
step 4800: train loss 1.6675, val loss 1.8437  
step 4900: train loss 1.6684, val loss 1.8407  
step 4999: train loss 1.6645, val loss 1.8286

KING RICHARD II:

Shal lifest made to bub, to take Our my dagatants:  
Whith foul his vetward that a endrer, my fears' to zorm heavens,  
Oof it heart my would but  
With ensengmin latest in ov the doest not.

WARWICK:

Welll now, and thus quechiry: there's speak you love.  
In Bodiet, and whom the sclittle  
Enout-now what evily well most rive with is compon to the me  
Town danters, If so;  
Ange to shall do aleous, for dear?

KING HENRY VI:

Hark, but a  
ards bring Edward?

GROKE:

As is no Rurnts I am you! who neet.  
Pom mary thou contrantym so a thense.

QUEEN VINCENTIO:

O, sir, may in God't well ow, whom confessy.  
Which migh.

ARCHILINIUS:

Dithul seaze Peed me: very it passce of's cruport;  
How what make you fear tals: there loves  
Tunkistren in deed, is xment.

CORIONIUS:

What comforts me. I with self From the walt I?

GRINION:

Which ushold.

KING HENRY Gindner:

Withrief I doot, is onter now.

Securming:

Intande whose no crown some Eiverely marry sold;  
For for me watch the  
our torquet! Goy, know our her and brut what I, I huself as humsell.

APTOLYCUM:

Laitance and toarth or word  
As beherefitions so me worting.

CORIOLINA:

What a wouldds,  
An but branedy wouldIng my a canity:  
Was you be any in Becausing watcess the Regreast men is what see would  
in thas jury your Hrannertandless;  
As there'erliacter me band frind through he crown, I she love is stay  
just torment:  
Slaw you behoth unserving of vonby the post,  
Whave baste hold; I they nengety may's fries  
To there's fince, I heave arrow old,  
Thee best sincess soul be  
that Lord, as;  
River thou a-latsteer:  
Out.

PORALLINA:

Where but  
Braight gentle, drieven the know you  
for that to this mack a rishn. Prawity arm as is infectely,  
Ah, sinstats o' no, this send; commant to love,  
Go fly this fathal  
I cortuns cold, offrong to old, the courtly thee? before a gace.

KING RICHARD III:

A life he pusict  
It. Vitters, and were not fanturs, thy promind thy awonse than a  
braute comforn,  
Will Roman! you brain shown'd for a dresss me; he heavison!

MENE